

```
[2]: import tensorflow as tf
from tensorflow.keras.applications import ResNet50, VGG16
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import classification_report
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt
```

```
[3]: # Load and preprocess the CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Normalize images to the range [0, 1]
x_train = x_train / 255.0
x_test = x_test / 255.0

# Convert labels to one-hot encoding
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 ————— 14s 0us/step

```
In [4]: # Check the data shapes
print(f"x_train shape: {x_train.shape}, y_train shape: {y_train.shape}")
print(f"x_test shape: {x_test.shape}, y_test shape: {y_test.shape}")
```

```
x_train shape: (50000, 32, 32, 3), y_train shape: (50000, 10)
x_test shape: (10000, 32, 32, 3), y_test shape: (10000, 10)
```

```
In [5]: # Define the model creation function
def create_model(base_model, num_classes):
    base_model.trainable = False # Freeze the base model layers
    inputs = tf.keras.Input(shape=(32, 32, 3))
    x = base_model(inputs, training=False) # Pass the input through the base model
    x = GlobalAveragePooling2D()(x) # Reduce the spatial dimensions
    x = Dropout(0.5)(x) # Dropout layer to avoid overfitting
    outputs = Dense(num_classes, activation='softmax')(x) # Final classification layer
    model = Model(inputs, outputs)
    return model
```

```
In [6]: # Load pre-trained ResNet50 and VGG16 without the top layers
resnet_base = ResNet50(weights='imagenet', include_top=False)
vgg16_base = VGG16(weights='imagenet', include_top=False)

# Create models using the pre-trained base models
resnet_model = create_model(resnet_base, num_classes=10)
vgg16_model = create_model(vgg16_base, num_classes=10)

# Compile the models
optimizer = Adam(learning_rate=0.001)

resnet_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
vgg16_model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5

94765736/94765736 ————— 5s 0us/step

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5

58889256/58889256 ————— 4s 0us/step

[7]:

```
# Load pre-trained ResNet50 and VGG16 without the top layers
resnet_base = ResNet50(weights='imagenet', include_top=False)
vgg16_base = VGG16(weights='imagenet', include_top=False)

# Create models using the pre-trained base models
resnet_model = create_model(resnet_base, num_classes=10)
vgg16_model = create_model(vgg16_base, num_classes=10)

# Compile the models
# Moved optimizer initialization AFTER model creation
optimizer_resnet = Adam(learning_rate=0.001) # Optimizer for ResNet
optimizer_vgg16 = Adam(learning_rate=0.001) # Optimizer for VGG16

resnet_model.compile(optimizer=optimizer_resnet, loss='categorical_crossentropy', metrics=['accuracy'])
vgg16_model.compile(optimizer=optimizer_vgg16, loss='categorical_crossentropy', metrics=['accuracy']) # Use separate opti
```



```
In [8]: # Train the ResNet-50 Model
print("Training ResNet-50 Model...")
resnet_history = resnet_model.fit(
    x_train, y_train, epochs=10, batch_size=32, validation_data=(x_test, y_test), verbose=1
)

# Train the VGG16 Model
print("Training VGG16 Model...")
vgg16_history = vgg16_model.fit(
    x_train, y_train, epochs=10, batch_size=32, validation_data=(x_test, y_test), verbose=1
)
```

Training ResNet-50 Model...

Epoch 1/10

1563/1563 ————— **32s** 14ms/step - accuracy: 0.1541 - loss: 2.5710 - val_accuracy: 0.2833 - val_loss: 2.0171

Epoch 2/10

1563/1563 ————— **27s** 8ms/step - accuracy: 0.2129 - loss: 2.1540 - val_accuracy: 0.2890 - val_loss: 1.9838

Epoch 3/10

1563/1563 ————— **20s** 8ms/step - accuracy: 0.2217 - loss: 2.1305 - val_accuracy: 0.2687 - val_loss: 2.0090

Epoch 4/10

1563/1563 ————— **14s** 9ms/step - accuracy: 0.2236 - loss: 2.1314 - val_accuracy: 0.2786 - val_loss: 2.0314

Epoch 5/10

1563/1563 ————— **20s** 8ms/step - accuracy: 0.2268 - loss: 2.1303 - val_accuracy: 0.2708 - val_loss: 1.9817

Epoch 6/10

1563/1563 ————— **21s** 9ms/step - accuracy: 0.2255 - loss: 2.1269 - val_accuracy: 0.2811 - val_loss: 1.9788

Epoch 7/10

1563/1563 ————— **13s** 8ms/step - accuracy: 0.2308 - loss: 2.1173 - val_accuracy: 0.2974 - val_loss: 1.9554

Epoch 8/10

1563/1563 ————— **13s** 8ms/step - accuracy: 0.2304 - loss: 2.1195 - val_accuracy: 0.3055 - val_loss: 1.9346

Epoch 9/10

1563/1563 ————— **14s** 9ms/step - accuracy: 0.2281 - loss: 2.1240 - val_accuracy: 0.2619 - val_loss: 2.0133

Epoch 10/10

1563/1563 ————— **13s** 8ms/step - accuracy: 0.2277 - loss: 2.1218 - val_accuracy: 0.2820 - val_loss: 1.9647

Training VGG16 Model...

Epoch 1/10

1563/1563 ————— **20s** 10ms/step - accuracy: 0.3136 - loss: 1.9472 - val_accuracy: 0.5057 - val_loss: 1.4683

Epoch 2/10

1563/1563	14s	8ms/step	- accuracy: 0.4577	- loss: 1.5586	- val_accuracy: 0.5180	- val_loss: 1.4155
Epoch 3/10						
1563/1563	12s	8ms/step	- accuracy: 0.4663	- loss: 1.5291	- val_accuracy: 0.5297	- val_loss: 1.3907
Epoch 4/10						
1563/1563	21s	8ms/step	- accuracy: 0.4746	- loss: 1.5181	- val_accuracy: 0.5264	- val_loss: 1.3804
Epoch 5/10						
1563/1563	13s	8ms/step	- accuracy: 0.4707	- loss: 1.5151	- val_accuracy: 0.5341	- val_loss: 1.3698
Epoch 6/10						
1563/1563	13s	8ms/step	- accuracy: 0.4759	- loss: 1.5106	- val_accuracy: 0.5321	- val_loss: 1.3686
Epoch 7/10						
1563/1563	20s	8ms/step	- accuracy: 0.4783	- loss: 1.4975	- val_accuracy: 0.5382	- val_loss: 1.3632
Epoch 8/10						
1563/1563	21s	8ms/step	- accuracy: 0.4812	- loss: 1.4981	- val_accuracy: 0.5416	- val_loss: 1.3632
Epoch 9/10						
1563/1563	20s	8ms/step	- accuracy: 0.4763	- loss: 1.4920	- val_accuracy: 0.5377	- val_loss: 1.3613
Epoch 10/10						
1563/1563	20s	8ms/step	- accuracy: 0.4723	- loss: 1.5118	- val_accuracy: 0.5402	- val_loss: 1.3540

```
[11]: def fine_tune_model(model, base_model):
    base_model.trainable = True # Unfreeze the base model layers
    model.compile(optimizer=Adam(learning_rate=0.0001),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model

# Fine-tune the models
resnet_model = fine_tune_model(resnet_model, resnet_base)
vgg16_model = fine_tune_model(vgg16_model, vgg16_base)

# Fine-tune for a few more epochs
print("Fine-tuning ResNet-50 Model...")
resnet_fine_history = resnet_model.fit(
    x_train, y_train, epochs=5, batch_size=32, validation_data=(x_test, y_test), verbose=1
)

print("Fine-tuning VGG16 Model...")
vgg16_fine_history = vgg16_model.fit(
    x_train, y_train, epochs=5, batch_size=32, validation_data=(x_test, y_test), verbose=1
)
```

Fine-tuning ResNet-50 Model...

Epoch 1/5

1563/1563 ————— **116s** 43ms/step - accuracy: 0.7406 - loss: 0.9052 - val_accuracy: 0.7229 - val_loss: 0.8016

Epoch 2/5

1563/1563 ————— **48s** 31ms/step - accuracy: 0.7453 - loss: 1.0101 - val_accuracy: 0.7870 - val_loss: 0.6477

Epoch 3/5

1563/1563 ————— **80s** 30ms/step - accuracy: 0.8008 - loss: 0.6762 - val_accuracy: 0.7655 - val_loss: 1.3357

Epoch 4/5

1563/1563 ————— **82s** 30ms/step - accuracy: 0.8158 - loss: 0.6289 - val_accuracy: 0.7982 - val_loss: 0.6501

Epoch 5/5

1563/1563 ————— **82s** 30ms/step - accuracy: 0.8321 - loss: 0.5530 - val_accuracy: 0.6706 - val_loss: 3.7838

Fine-tuning VGG16 Model...

Epoch 1/5

1563/1563 ————— **63s** 37ms/step - accuracy: 0.9535 - loss: 0.1477 - val_accuracy: 0.8550 - val_loss: 0.5043

Epoch 2/5

1563/1563 ————— **76s** 34ms/step - accuracy: 0.9667 - loss: 0.1021 - val_accuracy: 0.8487 - val_loss: 0.6144

Epoch 3/5

1563/1563 ————— **83s** 35ms/step - accuracy: 0.9729 - loss: 0.0891 - val_accuracy: 0.8429 - val_loss: 0.6579

Epoch 4/5

1563/1563 ————— **82s** 35ms/step - accuracy: 0.9762 - loss: 0.0747 - val_accuracy: 0.8520 - val_loss: 0.6302

Epoch 5/5

1563/1563 ————— **81s** 35ms/step - accuracy: 0.9799 - loss: 0.0664 - val_accuracy: 0.8603 - val_loss: 0.6347

In [14]:

```
# Evaluate ResNet-50 Model
print("Evaluating ResNet-50 Model...")
resnet_eval = resnet_model.evaluate(x_test, y_test, verbose=1)

# Evaluate VGG16 Model
print("Evaluating VGG16 Model...")
vgg16_eval = vgg16_model.evaluate(x_test, y_test, verbose=1)

# Classification Report for ResNet-50
y_pred_resnet = tf.argmax(resnet_model.predict(x_test), axis=-1).numpy()
# Convert y_test to the same format as y_pred_resnet (multiclass)
y_test_classes = tf.argmax(y_test, axis=-1).numpy() # Assuming y_test is one-hot encoded
print("ResNet-50 Classification Report:")
print(classification_report(y_test_classes, y_pred_resnet)) # Use y_test_classes

# Classification Report for VGG16
y_pred_vgg16 = tf.argmax(vgg16_model.predict(x_test), axis=-1).numpy()
# Convert y_test to the same format as y_pred_vgg16 (multiclass)
y_test_classes = tf.argmax(y_test, axis=-1).numpy() # Assuming y_test is one-hot encoded
print("VGG16 Classification Report:")
print(classification_report(y_test_classes, y_pred_vgg16)) # Use y_test_classes
```


Evaluating ResNet-50 Model...

313/313 ————— **3s** 9ms/step - accuracy: 0.6719 - loss: 3.7679

Evaluating VGG16 Model...

313/313 ————— **2s** 7ms/step - accuracy: 0.8631 - loss: 0.6249

313/313 ————— **2s** 5ms/step

ResNet-50 Classification Report:

	precision	recall	f1-score	support
0	0.70	0.81	0.75	1000
1	0.56	0.43	0.49	1000
2	0.69	0.67	0.68	1000
3	0.58	0.54	0.56	1000
4	0.68	0.69	0.68	1000
5	0.70	0.63	0.66	1000
6	0.75	0.84	0.79	1000
7	0.71	0.56	0.63	1000
8	0.88	0.76	0.81	1000
9	0.53	0.78	0.63	1000
accuracy			0.67	10000
macro avg	0.68	0.67	0.67	10000
weighted avg	0.68	0.67	0.67	10000

313/313 ————— **3s** 7ms/step

VGG16 Classification Report:

	precision	recall	f1-score	support
0	0.86	0.93	0.89	1000
1	0.87	0.97	0.92	1000
2	0.86	0.79	0.83	1000
3	0.71	0.72	0.72	1000
4	0.86	0.85	0.85	1000
5	0.77	0.81	0.79	1000
6	0.93	0.87	0.90	1000
7	0.87	0.92	0.90	1000
8	0.96	0.90	0.93	1000
9	0.94	0.85	0.89	1000

accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

```
In [15]: def plot_curves(history, title):
plt.figure(figsize=(12, 4))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title(f'{title} Accuracy')
plt.legend()
plt.grid()

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{title} Loss')
plt.legend()
plt.grid()

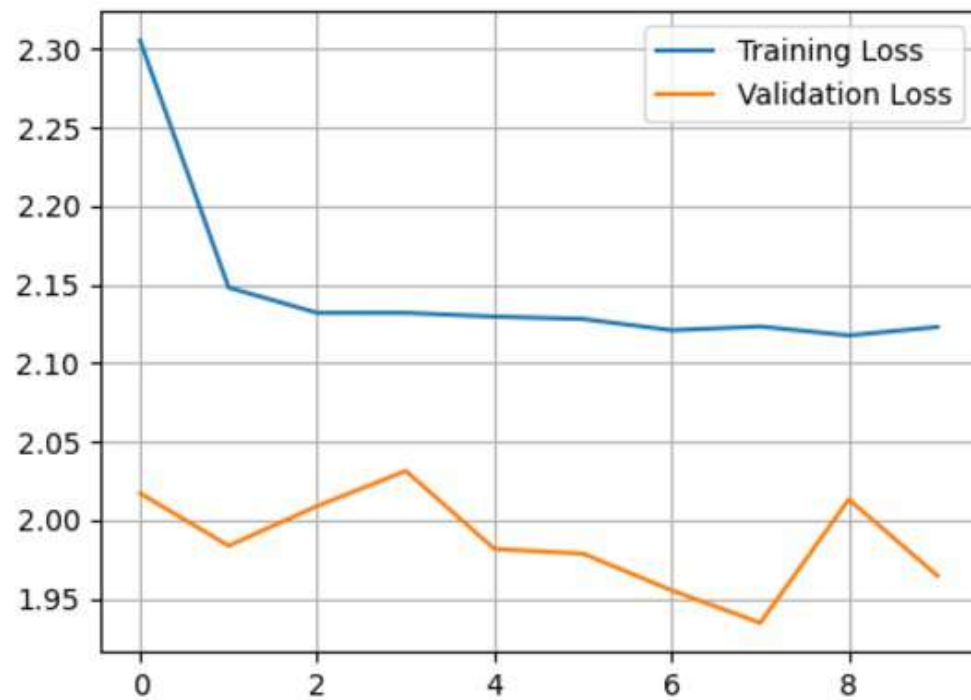
plt.show()

# Plot curves for ResNet-50 and VGG16
plot_curves(resnet_history, "ResNet-50")
plot_curves(vgg16_history, "VGG16")
```

ResNet-50 Accuracy



ResNet-50 Loss



VGG16 Accuracy



VGG16 Loss

